

Kaiwen Zhou

Updated January 4, 2025

Email: kzhou35@ucsc.edu

Github: [//github.com/KevinZ-01](https://github.com/KevinZ-01)

Webpage: <https://kevinz-01.github.io/>

Research interests

AI Agents, (Multimodal) Large Language Models, Responsible AI

Education

University of California, Santa Cruz

Ph.D. in Computer Science and Engineering

Sep. 2021 – Present

Advisor: Prof. Xin Eric Wang.

Zhejiang University

B.S. in Statistics

Sep. 2017 – June 2021

Work experience

Samsung Research America

Research intern

June 2024 – Sep. 2024

Mentor: Yilin Shen

Honda Research Institute

Research intern

April 2023 – Dec. 2023

Mentor: Kwonjoon Lee

Samsung Research America

Research intern

June 2022 – Sep. 2022

Mentor: Yilin Shen

Publications

Multimodal Situational Safety

Kaiwen Zhou, Chengzhi Liu, Xuandong Zhao, Anderson Compalas, Dawn Song, Xin Eric Wang

NeurIPS Workshop on RBFM 2024 *Oral*

Muffin or Chihuahua? Challenging Large Vision-Language Models with Multi-panel VQA

Yue Fan, Jing Gu, **Kaiwen Zhou**, Qianqi Yan, Shan Jiang, Ching-Chen Kuo, Xinze Guan, Xin Eric Wang

ACL 2024.

ViCor: Bridging Visual Understanding and Commonsense Reasoning with Large Language Models

Kaiwen Zhou, Kwonjoon Lee, Teruhisa Misu, Xin Eric Wang.

Findings of ACL 2024.

Navigation as the Attacker Wishes? Towards Building Byzantine-Robust Embodied Agents under Federated Learning

Yunchao Zhang, Zonglin Di, **Kaiwen Zhou**, Cihang Xie, Xin Eric Wang.

NAACL 2024

ESC: Exploration with Soft Commonsense Constraints for Zero-shot Object Navigation

Kaiwen Zhou, Kaizhi Zheng, Connor Pryor, Yilin Shen, Hongxia Jin,

Lise Getoor, Xin Eric Wang.

ICML 2023

FedVLN: Privacy-preserving Federated Vision-and-Language Navigation

Kaiwen Zhou, Xin Eric Wang.

ECCV 2022

JARVIS: A Neuro-Symbolic Commonsense Reasoning Framework for Conversational Embodied Agents

Kaizhi Zheng*, **Kaiwen Zhou***, Jing Gu*, Yue Fan*, Jialu Wang*, Zonglin Di, Xuehai He, Xin Eric Wang.

Socal NLP 2022

Selected research

Multimodal Situational Safety

Advisor: Prof. Xin Eric Wang

April. 2024 – Sep. 2024

- We propose a novel safety problem where the situation indicated by the visual input is a critical factor that influences the safety of the user's intent behind the query.
- We benchmark SOTA MLLMs and perform in-depth analysis.
- We propose multi-agent pipelines to improve situational safety performance.

Visual Commonsense Reasoning with LLM and VLMs

Advisor: Dr. Kwonjoon Lee, Prof. Xin Eric Wang

Mar. 2023 – Sep. 2023

- We studied the problem of visual commonsense reasoning and defined it into two sub-tasks: visual commonsense inference and visual commonsense understanding.
- We proposed a framework maximizing the capability of LLMs and VLMs to solve them.

LLM Commonsense Reasoning for Zero-shot Object Navigation

Advisor: Prof. Xin Eric Wang, Dr. Yilin Shen

June 2022 – Jan. 2023

- We proposed a framework that combines the commonsense reasoning of pre-trained LLM and classical navigation methods via Probabilistic Soft Logic (PSL) for zero-shot object navigation.
- We achieve SOTA zero-shot object navigation performance.

Privacy-preserving Federated Vision-and-Language Navigation

Advisor: Prof. Xin Eric Wang

Sep. 2021 – Mar. 2022

- We study the data privacy problem of VLN and propose a federated learning framework for vision and language navigation.
- We preserve the training and inference data privacy with comparable results with centralized training and achieve the best performance in pre-exploration.

Skills

Programming

Python, Pytorch.

Other experience

Amazon Alexa Prize SimBot Challenge

Advisor: Prof. Xin Eric Wang

Jan. 2022 – Apr. 2023

We investigated the problem of dialog-based embodied instruction following on TEACH benchmark and won the **first place** public challenge in the first phase. In the second phase, we are building an interactive embodied agent that can finish diverse tasks cooperating with human players. We won the **third place** in the second phase.