

Kaiwen Zhou

🌐 kevinz-01.github.io | ✉ kaiwenzhou111@gmail.com

SUMMARY

AI safety researcher focused on frontier AI risks: dangerous capability evaluation and misalignment monitoring, with prior work across red-teaming, safety post-training, and agent safety evaluation. First-author publications at ICML, ICLR, EMNLP, ACL, and ECCV; Area Chair for ARR.

PROFESSIONAL EXPERIENCE

Research Scientist, Scale AI (Preparedness)

June 2026 – Present

- Safety evaluation for recursive self-improvement (RSI) models.
- Misalignment evaluation for multi-agent systems.

Anthropic AI Safety Fellow (MATS) Mentor: William Saunders

Jan. 2026 – June 2026

- Defined a taxonomy of concepts indicating misaligned thought processes, and built a synthetic data generation pipeline for training concept monitors.
- Trained and validated linear probes that detect misaligned reasoning in frontier LLMs during multi-step conversations and tool use (*ICML 2026 Workshop*).

Research Intern, Microsoft Responsible AI Mentor: Ahmed Elgohary

Jun. 2025 – Sep. 2025

- Developed a red-teaming framework for LLM agents that iteratively crafts adversarial attacks.
- Built an effective and efficient red-teamer trained via distilled structured reasoning using SFT and RL.
- **Impact:** Deployed in Microsoft RAI product; a first-author paper (*Findings of EACL 2026*).

Research Intern, Samsung Research America Mentor: Yilin Shen

Jun. 2024 – Sep. 2024

- Developed prototype LLM-based agents for coding, scientific idea verification, and literature search.

Research Intern, Honda Research Institute Mentor: Kwonjoon Lee

Apr. 2023 – Dec. 2023

- Developed a novel framework for visual reasoning, maximizing the capability of foundation models.
- Achieved state-of-the-art training-free performance on visual reasoning tasks (*Findings of ACL 2024*).

Research Intern, Samsung Research America Mentor: Yilin Shen

Jun. 2022 – Sep. 2022

- Combined LLM reasoning with Probabilistic Soft Logic (PSL) for zero-shot object navigation.
- Achieved state-of-the-art performance in zero-shot embodied navigation tasks (*ICML 2023*).

EDUCATION

University of California, Santa Cruz

Sep. 2021 – June 2026

Ph.D. in Computer Science and Engineering

Advisor: Prof. Xin Eric Wang.

Research focus: AI safety, Alignment, (Embodied) AI agents.

Zhejiang University

Sep. 2017 – June 2021

B.S. in Statistics

SELECTED PUBLICATIONS

- **[ICML 2026 W]** Probing the Misaligned Thinking Process of Language Models
Kaiwen Zhou, Constantin Venhoff, Jonathan Michala, Xin Eric Wang, William Saunders
- **[In submission]** SafePro: Evaluating the Safety of Professional-Level AI Agents
Kaiwen Zhou, Shreedhar Jangam, Ashwin Nagarajan, Tejas Polu, Suhas Oruganti, ..., Xin Eric Wang.
- **[EACL 2026 Findings]** SIRAJ: Diverse and Efficient Red-Teaming for LLM Agents via Distilled Structured Reasoning
Kaiwen Zhou, Ahmed Elgohary, A S M Iftekhar, Amin Saied.

- [ICLR 2026] Presenting a Paper is an Art: Self-Improvement Aesthetic Agents for Academic Presentations
Chengzhi Liu*, Yuzhe Yang*, **Kaiwen Zhou**, Zhen Zhang, Yue Fan, Yannan Xie, Peng Qi, Xin Eric Wang.
- [EMNLP 2025] SafeKey: Amplifying Aha-Moment Insights for Safety Reasoning
Kaiwen Zhou, Xuandong Zhao, Gaowen Liu, Jayanth Srinivasa, Aosong Feng, Dawn Song, Xin Eric Wang.
- [IJCNLP-AACL 2025] The Hidden Risks of Large Reasoning Models: A Safety Assessment of R1
Kaiwen Zhou, Chengzhi Liu, Xuandong Zhao, Shreedhar Jangam, ..., Dawn Song, Xin Eric Wang.
- [ICLR 2025] Multimodal Situational Safety
Kaiwen Zhou*, Chengzhi Liu*, Xuandong Zhao, Anderson Compalas, Dawn Song, Xin Eric Wang.
- [ACL 2024] Muffin or Chihuahua? Challenging Large Vision-Language Models with Multipanel VQA
Yue Fan, Jing Gu, **Kaiwen Zhou**, Qianqi Yan, Shan Jiang, Ching-Chen Kuo, Xinze Guan, Xin Eric Wang.
- [ACL 2024 Findings] ViCor: Bridging Visual Understanding and Commonsense Reasoning with Large Language Models
Kaiwen Zhou, Kwonjoon Lee, Teruhisa Misu, Xin Eric Wang.
- [NAACL 2024] Navigation as the Attacker Wishes? Towards Building Byzantine-Robust Embodied Agents under Federated Learning
Yunchao Zhang, Zonglin Di, **Kaiwen Zhou**, Cihang Xie, Xin Eric Wang.
- [ICML 2023] ESC: Exploration with Soft Commonsense Constraints for Zero-shot Object Navigation
Kaiwen Zhou, Kaizhi Zheng, Connor Pryor, Yilin Shen, Hongxia Jin, Lise Getoor, Xin Eric Wang.
- [NeSy 2025 (Oral)] JARVIS: A Neuro-Symbolic Commonsense Reasoning Framework for Conversational Embodied Agents
Kaizhi Zheng*, **Kaiwen Zhou***, Jing Gu*, Yue Fan*, Jialu Wang*, Zonglin Di, Xuehai He, Xin Eric Wang.
- [ECCV 2022] FedVLN: Privacy-preserving Federated Vision-and-Language Navigation
Kaiwen Zhou, Xin Eric Wang.

SELECTED RESEARCH PROJECTS

AGI Safety: Safety Evaluation for Professional-Level AI Agents Oct. 2025 – Jan. 2026

Develop a safety evaluation dataset with safety risks in professional-level agentic tasks. Build an agent safety evaluation framework. Identify safety gaps of current AI models.

Improving the Safety Alignment of Large Reasoning Models March 2025 – May 2025

Identify the safety aha-moment of large reasoning models (LRMs), and amplify it for safer LRM with the proposed SafeKey training method, leading to significant safety improvement.

Multimodal Situational Safety Apr. 2024 – Sep. 2024

Propose a novel safety problem where the situation in visual input affects the safety of the user’s intent in chat and embodied scenarios; benchmark MLLMs and propose multi-agent pipelines to improve situational safety.

Amazon Alexa Prize SimBot Challenge Jan. 2022 – Apr. 2023

Build dialog-based embodied instruction following agent; won first place in the public challenge (phase I) and third place in real-user interaction stage (phase II).

AI TECHNICAL SKILLS

Red-teaming, safety evaluation, safety monitoring & probing, post-training, alignment, reinforcement learning, supervised fine-tuning, reasoning, multimodal LLMs